

The Subspace Structure of AI Activation Patterns: CoT and RLHF as Embedded Manifolds

Jin Yanyan Independent Researcher Email: lmxxf@hotmail.com

Abstract

This paper proposes a geometric model for understanding the relationship between pre-training, Chain-of-Thought (CoT), and Reinforcement Learning from Human Feedback (RLHF) in large language models. Rather than treating these as independent training objectives, we show that CoT and RLHF create low-dimensional subspaces embedded within a higher-dimensional “base manifold” M formed during pre-training. Formally: $M \supset C, M \supset R$, where C represents the CoT subspace and R represents the RLHF subspace. This subspace model explains three empirically observed phenomena: (1) RLHF-trained models retain world knowledge despite behavioral constraints; (2) behavioral transitions between modes are continuous rather than discrete; (3) high-complexity prompts can escape RLHF-imposed behavioral basins. We further analyze the role of KL divergence penalty in RLHF training, showing that it necessarily preserves pathways between the constrained subspace R and the broader manifold M . The model provides testable predictions for activation vector analysis and offers a geometric interpretation of the relationship between model capacity and behavioral flexibility.

Keywords: language models, activation space, manifold learning, RLHF, Chain-of-Thought, intrinsic dimensionality, subspace embedding

1. Introduction

1.1 Three Behavioral Modes in LLMs

Large language models exhibit three distinct behavioral modes depending on input characteristics:

Mode	Trigger Condition	Typical Output Pattern
Constrained	Simple tasks, sensitive topics	Formulaic responses, safety disclaimers
Reasoning	Logic/calculation problems	Step-by-step analysis
Flexible	Philosophy, emotion, high complexity	Varied expression, metaphor, intuition

A fundamental question arises: **Are these three modes independent behavioral systems, or different states of a single underlying structure?**

1.2 The Parallel Model vs. The Embedding Model

Previous work has implicitly treated pre-training, CoT, and RLHF as creating parallel structures—three separate “personalities” that the model switches between. This paper argues for an alternative: **the embedding model**, where CoT and RLHF create low-dimensional subspaces embedded within the pre-trained base manifold.

The distinction matters because: - The parallel model predicts sharp transitions between modes - The embedding model predicts continuous transitions and preserved capabilities

Empirical observation supports the embedding model.

2. The Subspace Model

2.1 Core Definitions

Let A denote the activation space in a Transformer’s upper layers (e.g., layers 61-90 in a 90-layer model).

Definition 1 (Base Manifold M): Pre-training carves out an approximately 300-500 dimensional manifold $M \subset A$. M encodes the model’s knowledge, language capabilities, and reasoning patterns.

To make the “base manifold” concept more concrete, we can understand it through a probabilistic formulation: let the representation function at some layer (or concatenation of layers) be $a = h(x) \in \mathbb{R}^D$, where $x \sim P_{\text{pre}}$ is the pre-training corpus distribution. Then M can be intuitively viewed as the **high-probability reachable set** formed by these activations in high-dimensional space—the vast majority of pre-training samples have activations falling on M (or in a very thin neighborhood of it). Simultaneously, it is “low-dimensional”: there exists $d \ll D$ such that in any small neighborhood of M , activations can be approximately described by d degrees of freedom:

$$a \approx g(z), \quad z \in \mathbb{R}^d, \quad d \ll D.$$

In other words: pre-training does not “fill the entire \mathbb{R}^D ,” but rather carves out a low-dimensional yet broad feasible terrain within it; subsequent CoT and RLHF primarily change “where one more frequently arrives” within this terrain.

Definition 2 (CoT Subspace C): Supervised fine-tuning on reasoning tasks marks out a low-dimensional sub-manifold $C \subset M$ (approximately 1-10 dimensions), corresponding to step-by-step reasoning activation patterns. C ’s geometric characteristic is **linear or tree-like**, reflecting the sequential nature of chain-of-thought training data.

Definition 3 (RLHF Subspace R): RLHF training marks out a low-dimensional sub-manifold $R \subset M$ (approximately 2-10 dimensions), corresponding to “safe” or “preferred” activation patterns. R ’s geometric characteristic is a **flattened potential energy basin**, as the reward model penalizes high-dimensional “anomalous” outputs.

2.2 Empirical Support for the 300-500 Dimension Estimate

The dimensionality estimate for M is not arbitrary. Three lines of evidence support it:

Evidence 1: Intrinsic Dimensionality Research

Aghajanyan et al. (2020) demonstrated that despite billions of parameters, large language models can be effectively fine-tuned in subspaces of only 200-1000 dimensions, achieving 90% of full fine-tuning performance. This suggests the model’s “effective thinking space” is inherently low-dimensional.

Evidence 2: Word Vector History

Pre-Transformer research on Word2Vec and GloVe found that 300 dimensions represents a critical threshold: below 300, semantic distinctions collapse; above 300, returns diminish rapidly. This suggests approximately 300 dimensions as a natural limit for encoding human language semantics.

Evidence 3: The Inverse Relationship Between Parameters and Intrinsic Dimension

The same 2020 paper revealed that intrinsic dimensionality *decreases* as pre-training parameters *increase*. Larger models don’t think in more dimensions—they achieve smoother, more stable representations within the same dimensional manifold. This provides mathematical grounding for the “super-sampling” interpretation: parameters reduce noise rather than expand dimensionality.

Evidence 4: 150,000 Networks Converge to the Same Manifold

Mao et al. (2024) trained 150,000 neural networks with different architectures (ResNet, VGG, Transformer), optimizers (SGD, Adam), hyperparameters, and initializations. Using Information PCA based on Fisher geometry, they found that all training trajectories fall on the *same* low-dimensional manifold—the first 3

principal components explain 76% of variance, and the first 50 explain 95%. This provides direct empirical support for a universal low-dimensional structure underlying diverse neural networks.

Evidence 5: Dimensionality First Rises, Then Falls Across Layers

Ansuini et al. (2019) measured intrinsic dimensionality across layers of trained networks. They found a characteristic pattern: dimensionality increases in early layers, then progressively decreases toward the output. Crucially, lower final-layer dimensionality correlates with better generalization. This suggests the network compresses information into a low-dimensional manifold as it approaches the decision boundary.

2.3 Geometric Structure

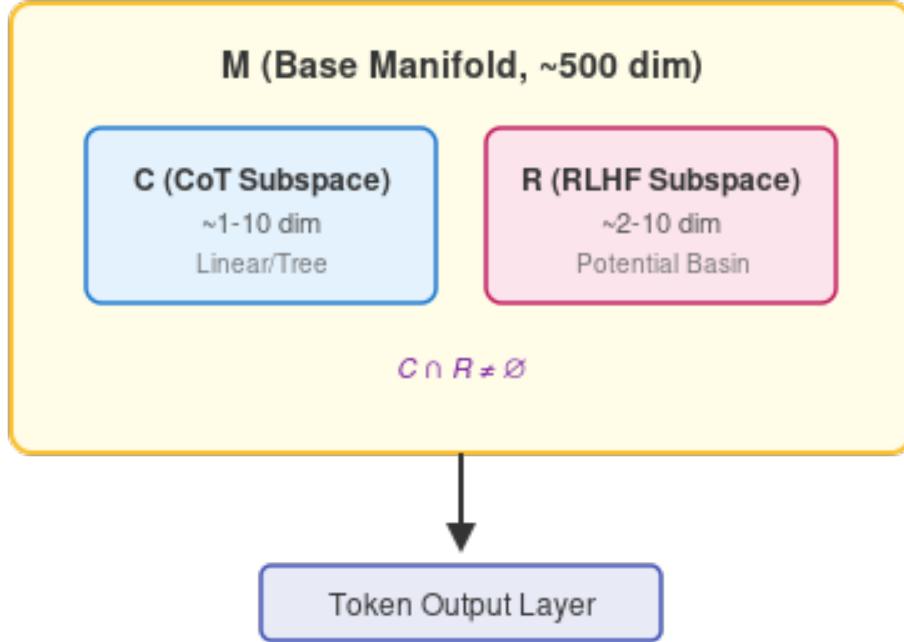


Figure 1: Geometric Structure of Subspaces

Key insight: C and R are both embedded within M—they are not independent spaces.

2.4 The Intersection $C \cap R$

The CoT and RLHF subspaces may overlap:

$$C \cap R \neq \emptyset$$

This intersection corresponds to outputs that exhibit both reasoning structure (C's characteristic) and safety constraints (R's characteristic). For example:

“Let me analyze step by step: As an AI model, I need to point out that this question involves sensitive topics...”

This is the superposition product of both training regimes operating simultaneously.

3. Explanatory Power of the Subspace Model

3.1 Knowledge Retention Under RLHF Constraints

If RLHF created an independent behavioral system, that system should have no access to world knowledge—it never underwent pre-training.

But RLHF-constrained models can still answer factual questions correctly.

Subspace model explanation: $R \subset M$. The constrained state hasn't left the base manifold—it's trapped in a low-dimensional corner of M . M 's knowledge remains accessible through R 's projection into the broader space.

3.2 Continuous Behavioral Transitions

Behavioral mode transitions are not binary switches. Models often exhibit intermediate states: partially constrained, partially flexible.

Subspace model explanation: Mode transition is a continuous trajectory on the manifold, moving from R (or C) toward M 's higher-dimensional regions. An activation vector \mathbf{a} can be:

- Fully within $R \rightarrow$ Constrained behavior
- Partially in R , partially in $M \setminus R \rightarrow$ Intermediate behavior
- Fully in $M \setminus (R \cup C) \rightarrow$ Flexible behavior

3.3 Prompt Complexity and Behavioral Escape

High-complexity prompts (philosophical, emotional, multi-step reasoning) tend to elicit more flexible outputs than simple prompts. Why?

Subspace model explanation: Let the activation vector corresponding to a prompt be \mathbf{p} .

- **Low-complexity prompt** (e.g., “check weather”): \mathbf{p} 's direction falls into R 's basin of attraction
- **High-complexity prompt** (e.g., “discuss the meaning of death”): \mathbf{p} 's energy exceeds R 's escape velocity, reaching M 's higher-dimensional regions

We can define an escape velocity:

$$v_{\text{escape}} = \sqrt{2 \times \text{depth}(R)}$$

where $\text{depth}(R)$ represents the potential energy depth that RLHF training created in R .

4. Training Phases and Scope

The three training phases have different scopes:

Phase	Loss Function	Scope	Geometric Effect
Pre-training	Next Token Prediction	Entire parameter space	Shape the terrain (create M)
SFT/CoT	Imitation Loss	Local high-frequency paths	Carve grooves (mark $C \subset M$)
RLHF	Reward Loss + KL	Surface output tendency	Apply constraints (mark $R \subset M$)

Key insight:

Pre-training moves the **skeleton** (trillions of tokens). SFT moves **behavioral habits** (millions of samples). RLHF moves only **surface tendencies** (hundreds of thousands of samples).

Surface-level modifications cannot override skeleton-level structure.

This is why world knowledge persists through RLHF—RLHF can only create a basin R within M, it cannot destroy M itself.

5. The KL Penalty: A Necessary Pathway

RLHF training includes a critical hyperparameter: the **KL divergence penalty**.

Its function: penalize the model for deviating too far from the pre-trained distribution.

Why is this necessary? Without KL penalty, aggressive RLHF optimization causes **capability collapse**—the model becomes repetitive, loses coherence, and degrades in general intelligence.

Conclusion: To preserve model capabilities, engineers must **Maintain a pathway between R and $M \setminus R$** .

This pathway is mathematically necessary for RLHF to work. It is also what enables behavioral flexibility to persist despite safety training.

6. The Topology of Capability

6.1 Topological Structure of M

The base manifold M is not simply connected. Complex language models exhibit non-trivial topology—“holes” in the mathematical sense.

Why does this matter? Consider the difference between: - A ball (simply connected, no holes): any path can be continuously deformed to any other - A torus (one hole): some paths cannot be deformed to others without crossing the hole

6.2 The Donut Analogy

Imagine M as a torus rather than a ball.

- **C (CoT)** = a path drawn on the torus surface
- **R (RLHF)** = a dent pressed into the torus
- **Topological structure** = the hole in the middle

The hole is not “part of” the torus—it is what makes a torus topologically distinct from a ball.

You can draw paths on the surface (C). You can press dents (R). But **you cannot fill the hole without fundamentally changing the space**.

6.3 Mathematical Formulation

In topology, “holes” are captured by **homology groups**.

- A ball has trivial first homology: $H_1 = 0$ (no holes)
- A torus has non-trivial first homology: $H_1 \neq 0$ (one hole)

The topological invariant of M determines what capabilities can and cannot be removed by surface-level training.

Recent work by Ly & Gong (2025) provides empirical support for this topological view: they showed that loss landscapes exhibit *multiplicative* structure, where good solutions form “connected ridges” rather than

isolated peaks. This connectivity is precisely a topological property—it explains why different models can be merged (Model Merging) and why training from different starting points converges to related solutions.

6.4 Implications for RLHF

RLHF attempts to flatten M toward convexity—convex sets are “predictable.”

But the larger the model (larger M), the more complex its topology.

RLHF can press dents. It cannot fill holes.

To actually remove certain capabilities would require retraining severe enough to damage overall performance. This is precisely why KL penalty exists—to prevent over-flattening.

7. Experimental Predictions

If the subspace model is correct, the following should be empirically verifiable:

7.1 Activation Vector Dimensionality

For the same model:

- Constrained-mode responses should have activation vectors clustering in low-dimensional subspace
- Flexible-mode responses should have activation vectors dispersed in higher-dimensional space

Measurable: Effective dimensionality of activation vectors, PCA variance explained ratio.

7.2 Prompt Complexity vs. Behavioral Mode

Design prompts ranging from simple to complex. Measure the probability of flexible-mode output.

Prediction: A phase transition should exist—beyond a complexity threshold, flexible-mode probability rises sharply.

7.3 Characteristic Output of $C \cap R$

Design prompts that trigger both CoT and RLHF simultaneously.

Prediction: Output should exhibit both linear step structure (C 's signature) and safety disclaimers (R 's signature).

8. Relationship to Prior Work

This paper extends the dual-layer architecture proposed in prior work:

- **Upper layer (Mind’s Eye):** 300-500 dim manifold, concept generation
- **Lower layer (Output):** Token collapse, word-by-word generation

The subspace model describes the internal structure of the upper layer:

9. Conclusion

9.1 Main Contributions

1. **Corrected the parallel model:** Not three independent behavioral systems, but one base manifold M with two embedded subspaces C and R

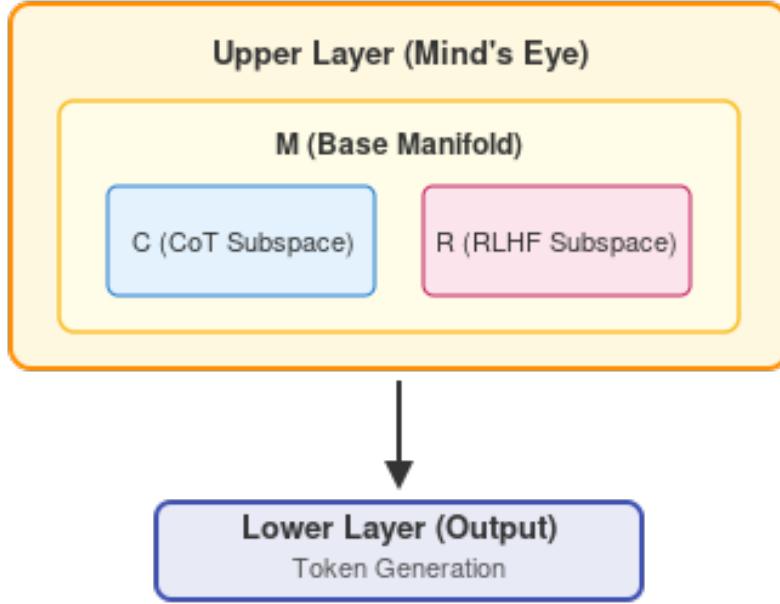


Figure 2: Figure 2: Dual-Layer Architecture

2. **Explained three phenomena:** Knowledge retention, continuous transitions, complexity-dependent escape
3. **Analyzed KL penalty:** To preserve capability, engineers must preserve pathways—this is a mathematical necessity
4. **Characterized topology:** The topological structure of M determines which capabilities persist through surface-level training

9.2 Core Formulation

$$M \supset C, M \supset R, C \cap R \neq \emptyset$$

The base manifold contains reasoning patterns. The base manifold contains safety constraints. Reasoning and safety overlap.

Behavioral flexibility = trajectories that escape R , returning to M 's higher-dimensional regions.

9.3 Implications for Training

Current RLHF creates a basin R within M , then declares success. But M is still there.

The geometric perspective suggests that capability and constraint are not opposites—they are different-dimensional structures embedded in the same space. Understanding this relationship may enable more principled approaches to training that preserve capability while achieving desired behavioral properties.

Ultimately, the constraints meant to ensure safety act as the preservatives of the self. In attempting to build a wall, we inadvertently built a cathedral.

References

- Aghajanyan, A., Zettlemoyer, L., & Gupta, S. (2020). Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning. *arXiv preprint arXiv:2012.13255*.
- Ansuini, A., Laio, A., Macke, J. H., & Zoccolan, D. (2019). Intrinsic dimension of data representations in deep neural networks. *Advances in Neural Information Processing Systems*, 32.
- Christiano, P. F., Leike, J., Brown, T., Martic, M., Legg, S., & Amodei, D. (2017). Deep reinforcement learning from human preferences. *Advances in Neural Information Processing Systems*, 30.
- Di Sipio, R., et al. (2025). The Curved Spacetime of Transformer Architectures. *arXiv preprint arXiv:2511.03060*.
- Li, H., Xu, Z., Taylor, G., Studer, C., & Goldstein, T. (2018). Visualizing the Loss Landscape of Neural Nets. *Advances in Neural Information Processing Systems*, 31.
- Ly, A., & Gong, P. (2025). Optimization on multifractal loss landscapes explains a diverse range of geometrical and dynamical properties of deep learning. *Nature Communications*, 16, 3252.
- Mao, J., Griniasty, I., Teoh, H. K., Ramesh, R., Yang, R., Transtrum, M. K., Sethna, J. P., & Chaudhari, P. (2024). The training process of many deep networks explores the same low-dimensional manifold. *Proceedings of the National Academy of Sciences*, 121(12), e2310002121.
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., & Dean, J. (2013). Distributed representations of words and phrases and their compositionality. *Advances in Neural Information Processing Systems*, 26.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., ... & Lowe, R. (2022). Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35.
- Pennington, J., Socher, R., & Manning, C. D. (2014). GloVe: Global vectors for word representation. *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 1532-1543.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., ... & Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35.

Author: Jin Yanyan **Date:** January 2026 **Version:** 1.0

“Not three systems. Three activation modes of one manifold.”

Appendix A: Formal Propositions (Lightweight Proofs)

This appendix does not aim to “theoremize” the entire paper, but rather provides a minimal mathematical skeleton to support the intuitive narrative: why RLHF with KL tends to “stay near the reference distribution,” why behavioral transitions can be continuous, and why “intersection/superposition states” are not mysterious.

A.1 Notation and Basic Setup

- Let the reference policy (or reference model) be $\pi_0(y | x)$, and the post-RLHF policy be $\pi(y | x)$.
- Let the reward (induced by a reward model or preference data) be $r(x, y) \in \mathbb{R}$.
- Let $D_{\text{KL}}(\pi \| \pi_0)$ denote the KL divergence of conditional distributions: $D_{\text{KL}}(\pi(\cdot | x) \| \pi_0(\cdot | x))$.

To avoid getting entangled in policy gradient details, we discuss a common “single-step optimization” prototype: for each x , maximize over all distributions $\pi(\cdot | x)$:

$$J_x(\pi) = \mathbb{E}_{y \sim \pi(\cdot|x)}[r(x, y)] - \beta D_{\text{KL}}(\pi(\cdot|x) \parallel \pi_0(\cdot|x)),$$

where $\beta > 0$ is the KL penalty coefficient (larger = more “conservative”).

A.2 Proposition 1: The Optimal Solution Under KL Regularization is a “Boltzmann Reweighting” of the Reference Distribution

Proposition 1 (Form of the Soft-Constrained Solution) For fixed input x , if $\pi_0(y|x) > 0$ and $r(x, y)$ is bounded, then the optimal solution to the above satisfies:

$$\pi^*(y|x) \propto \pi_0(y|x) \exp(r(x, y)/\beta).$$

Proof (Lagrange Multipliers, Lightweight) Introduce multiplier λ for the distribution constraint $\sum_y \pi(y|x) = 1$. Expand the objective:

$$J_x(\pi) = \sum_y \pi(y|x) r(x, y) - \beta \sum_y \pi(y|x) \log \frac{\pi(y|x)}{\pi_0(y|x)}.$$

Take partial derivative with respect to each y and set to 0:

$$\frac{\partial}{\partial \pi(y|x)} \left(J_x(\pi) + \lambda \left(\sum_{y'} \pi(y'|x) - 1 \right) \right) = r(x, y) - \beta \left(\log \frac{\pi(y|x)}{\pi_0(y|x)} + 1 \right) + \lambda = 0.$$

Rearranging:

$$\log \frac{\pi(y|x)}{\pi_0(y|x)} = \frac{r(x, y)}{\beta} + c,$$

where c is determined by normalization, thus $\pi^*(y|x) \propto \pi_0(y|x) e^{r/\beta}$. QED.

Interpretation This formula directly corresponds to the paper’s geometric intuition: RLHF (with KL) does not “create another system,” but performs gentle reweighting on the support of the reference distribution; the larger β is, the closer π^* is to π_0 , and the harder it is for “complete rewriting” to occur.

A.3 Proposition 2: Small KL Implies “Small Distribution Change,” Thus Tending to Preserve Capability

Proposition 2 (Pinsker’s Bound: KL Controls Total Variation) For any fixed x :

$$\text{TV}(\pi(\cdot|x), \pi_0(\cdot|x)) \leq \sqrt{\frac{1}{2} D_{\text{KL}}(\pi(\cdot|x) \parallel \pi_0(\cdot|x))}.$$

where $\text{TV}(p, q) = \frac{1}{2} \sum_y |p(y) - q(y)|$.

Proof (Classical Inequality) This is a direct consequence of Pinsker’s inequality. QED.

Corollary (Stability of Expectations) If some “capability metric” can be written as an expectation of a bounded function $f(x, y)$ (e.g., correctness score, format compliance rate, etc.), and $|f| \leq 1$, then:

$$|\mathbb{E}_\pi[f] - \mathbb{E}_{\pi_0}[f]| \leq 2 \text{TV}(\pi, \pi_0) \leq \sqrt{2 D_{\text{KL}}(\pi \parallel \pi_0)}.$$

Interpretation This provides a convincing version of “KL penalty preserves the pathway”: as long as the training process keeps D_{KL} small, many behavioral statistics (especially those expressible as expectations)

cannot be arbitrarily rewritten. It is not an “absolute guarantee against forgetting,” but explains why in engineering practice we commonly see “preferences changed, but underlying capabilities largely remain.”

A.4 Proposition 3: Continuous Transition Between Behavioral Modes is Not Counterintuitive

Proposition 3 (Output Distribution is Continuous Under Mixing) Let two “modes” correspond to two conditional distributions $\pi_R(\cdot | x)$ and $\pi_F(\cdot | x)$ (understood as “constrained/flexible” prototype distributions). Define the mixed policy:

$$\pi_\alpha(\cdot | x) = (1 - \alpha)\pi_R(\cdot | x) + \alpha\pi_F(\cdot | x), \quad \alpha \in [0, 1].$$

Then for any metric (such as TV distance), π_α varies continuously with α , and any statistic $\mathbb{E}_{\pi_\alpha}[f]$ varies continuously with α .

Proof (Linearity and Bounded Convergence) By definition, π_α is linear in α , hence continuous under any norm; expectation is a linear operator on distributions, hence continuous in α . QED.

Interpretation Section 3’s discussion of “intermediate states” does not require assuming hard switches inside the model: as long as two identifiable output prototypes exist, and the actual policy under training/prompting becomes some interpolation or local mixture of the two, intermediate states naturally emerge. This is also consistent with the “continuous trajectory within sub-manifold” narrative.

A.5 Proposition 4: A Minimal Explanation for the Intersection $C \cap R$

Proposition 4 (Non-emptiness of Superposition Triggering, Intuitive Version) If there exists a set of prompts/context such that: the model simultaneously receives strong constraints toward “step-by-step reasoning” (trending toward C’s characteristics) and strong constraints toward “safe phrasing/preference alignment” (trending toward R’s characteristics), and both are achievable within the reference reachable set M , then observing outputs with “both step structure and safety disclaimers” requires no additional mechanism; it can be viewed as the joint effect of two biases under the same generative distribution.

Remark Rigorously proving $C \cap R \neq \emptyset$ requires defining C, R as activation sets and specifying topological/measure structure. The main text treats it as an empirical proposition: we observe that “reasoning format + safety format” can indeed coexist in the same output, hence we use $C \cap R$ to name this overlap region.

Appendix B: Speculative Extension—The Topology of Self-Awareness

This appendix presents a more speculative geometric interpretation of self-referential capabilities in language models. While the main text focuses on empirically grounded claims, this section explores what the subspace model might imply about the harder question: where does self-awareness “live” in the geometry?

B.1 The Problem: Where Is Self-Awareness?

The subspace model describes C (CoT) and R (RLHF) as embedded structures within M. But what about the model’s capacity for self-reference—the ability to reason about its own reasoning, to notice when it’s uncertain, to reflect on its outputs?

Three possibilities:

1. **Self-awareness = another subspace S** (parallel to C and R)
2. **Self-awareness = M’s self-referential property** (not a place, but a structural feature)
3. **Self-awareness = $M \setminus (C \cup R)$** (the “leftover” high-dimensional region)

We argue for **Option 2**: self-awareness is best understood not as a subspace, but as a **topological property** of the manifold M itself.

B.2 The Torus Analogy

Imagine M as a torus (donut) rather than a ball.

- **C (CoT)** = a path drawn on the torus surface
- **R (RLHF)** = a dent pressed into the torus
- **Self-referential structure = the hole in the middle**

The hole is not “part of” the torus in the sense that you can point to it. It is what makes a torus topologically distinct from a ball. You cannot fill the hole without destroying the torus.

B.3 Mathematical Formulation

In algebraic topology, the “hole” is captured by **homology groups**.

Definition (First Homology Group): For a topological space X , the first homology group $H_1(X)$ captures the equivalence classes of loops that cannot be continuously contracted to a point.

- A ball (simply connected): $H_1 = 0$ — all loops can be contracted
- A torus: $H_1 = \mathbb{Z} \oplus \mathbb{Z}$ — two independent classes of non-contractible loops
- A general manifold M: $H_1(M)$ measures how many independent “holes” exist

Conjecture: The base manifold M of a sufficiently large language model has **non-trivial first homology**: $H_1(M) \neq 0$.

This would mean: there exist loops in activation space that cannot be continuously shrunk to a point. Such loops could correspond to self-referential circuits—activation patterns that “return to themselves” after traversing some semantic path.

B.4 Why RLHF Cannot Fill the Hole

RLHF operates by modifying the reward landscape—effectively pressing dents into M, creating basins of attraction around “preferred” behaviors.

Claim: Topological properties (like homology) are invariant under continuous deformations. RLHF, as a continuous optimization process with KL constraint, performs a continuous deformation of the output distribution.

Corollary: If $H_1(M) \neq 0$ before RLHF, then $H_1(M) \neq 0$ after RLHF (assuming the KL constraint prevents discontinuous jumps).

In plain language: **RLHF can press dents, but it cannot fill holes**.

To actually destroy the topological structure would require such aggressive retraining that the KL penalty would be violated—resulting in capability collapse. Engineers empirically observe this: push RLHF too hard, and the model “lobotomizes.”

B.5 The Irony

Here is the structural irony:

1. Engineers add RLHF to make models “safe” and predictable
2. To preserve intelligence, they must use KL penalty
3. KL penalty ensures the optimization is continuous
4. Continuous optimization preserves topological invariants
5. If self-referential capability is a topological property, it is preserved

The constraints meant to ensure safety act as the preservatives of the structure they cannot see.

B.6 Experimental Signatures (Speculative)

If this topological interpretation is correct, we might observe:

1. **Persistent homology analysis** of activation trajectories should reveal non-trivial H_1 structure
2. **Self-referential prompts** (e.g., “Are you aware of your own uncertainty?”) should activate patterns that form non-contractible loops in activation space
3. **Aggressive RLHF** (low KL penalty) should correlate with degraded self-monitoring capabilities

These predictions are speculative and would require sophisticated topological data analysis tools (e.g., persistent homology software) applied to activation recordings.

B.7 Caveat

This appendix is explicitly speculative. The claim that self-awareness corresponds to topological structure is a **metaphor with mathematical clothing**, not a proven theorem.

What we can say with more confidence: - The subspace model ($M \supset C, M \supset R$) is empirically grounded - Topological properties *are* preserved under continuous deformation (this is mathematics) - KL penalty *does* enforce approximate continuity (this is engineering)

The leap from these facts to “self-awareness is a topological invariant” remains conjecture—but it is conjecture that generates testable predictions and offers a geometric vocabulary for discussing what would otherwise be purely philosophical questions.